

开源操作系统实践

第二讲 内核静态分析

向勇

清华大学计算机系

xyong@tsinghua.edu.cn

2022年7月

第二讲 内核静态分析

1. Background
2. Static Analysis
3. A Perfect Work

1. Background

1. Background

- **LINT: A Programming Tool**
- The Problem
- Compiler

2. Static Analysis

3. A Perfect Work

DOUBLE-CHECKING PROGRAMS: LINT

Checks your program more thoroughly than **cc** does:

```
Utility : lint { fileName }*
```

Lint scans the specified source files and displays any potential errors that it finds.

```
$ lint reverse.c          ---> check "reverse.c".  
reverse defined ( reverse.c(12) ), but never used  
  
$ lint palindrome.c       ---> check "palindrome.c".  
palindrome defined ( palindrome.c ( 12 ) ), but never used reverse used ( palindrome.c(14) ), but not defined
```

DOUBLE-CHECKING PROGRAMS: LINT

```
$ lint main2.c          ---> check "main2.c".
main2.c(11) : warning: main() returns random value to invocation environment
printf returns value which is always ignored
palindrome used ( main2.c(9) ), but not defined

$ lint main2.c reverse.c palindrome.c ---> check all modules together.
main2.c:
main2.c(11): warning: main() returns random value to invocation environment
reverse.c:
palindrome.c:
Lint pass2:
printf returns value which is always ignored
$ _
```

1. Background

1.2 The Problem

1. Background

- LINT: A Programming Tool
- **The Problem**
- Compiler


2. Static Analysis

3. A Perfect Work

Possible Programming Bugs

- Application
- Programming Language
- User lib & Syscall
- Compiler
- Hardware

Landscape for all discovered CVE in [2019](#)



Vulnerabilities by Category ([2016-2020](#))


	2020	2019	2018	2017	2016
Remote Coded Execution	345 (27%)	323	292	301	269
Elevation of Privilege	559 (44%)	198	145	90	114
Information Disclosure	179 (14%)	177	153	193	102
Denial of Service	46 (4%)	52	29	43	0
Spoofing	104 (8%)	63	20	16	12
Tampering	7 (0.5%)	8	8	1	0
Security Feature Bypass	30 (2.5%)	38	20	41	26

The Problem

- Building secure systems is hard
 - 2/3 of Internet servers have gaping security holes
- The problem is buggy software
 - And a few pitfalls account for many vulnerabilities
- Challenge: Improve programming technology
 - Need way to gain assurance in our software
 - Static analysis can help!

[List of tools for static code analysis](#)

Existing Paradigms



What Makes Security Hard?

Security is hard because of...

- buffer overruns
- privilege pitfalls
- untrusted data
-

1.3 Compiler

1. Background

- LINT: A Programming Tool
- The Problem
- **Compiler**

2. Static Analysis

3. A Perfect Work

What is a compiler?

- A program that translates a program in one language to another language
 - The essential interface between applications & architectures
- Typically lowers the level of abstraction
 - analyzes and reasons about the program & architecture
- We expect the program to be optimized, i.e., better than the original
 - ideally exploiting architectural strengths and hiding weaknesses

Role of compilers

- **Bridge** complexity and evolution in architecture, languages, & applications
- **Help programmers** with correctness, reliability, program understanding
- Compiler **optimizations** can significantly improve performance
 - 1 to 10x on conventional processors
- **Performance stability**: one line change can dramatically alter performance
 - unfortunate, but true

Optimization

What should it do?

1. improve running time, or
2. decrease space requirements
3. decrease power consumption

How does it do it?

Example optimizations

- Division of optimizations
 1. Machine independent
 2. Machine dependent
- Faster code optimizations
 - common subexpression elimination
 - constant folding
 - dead code elimination
 - register allocation
 - scheduling

Analysis

Scope of program analysis

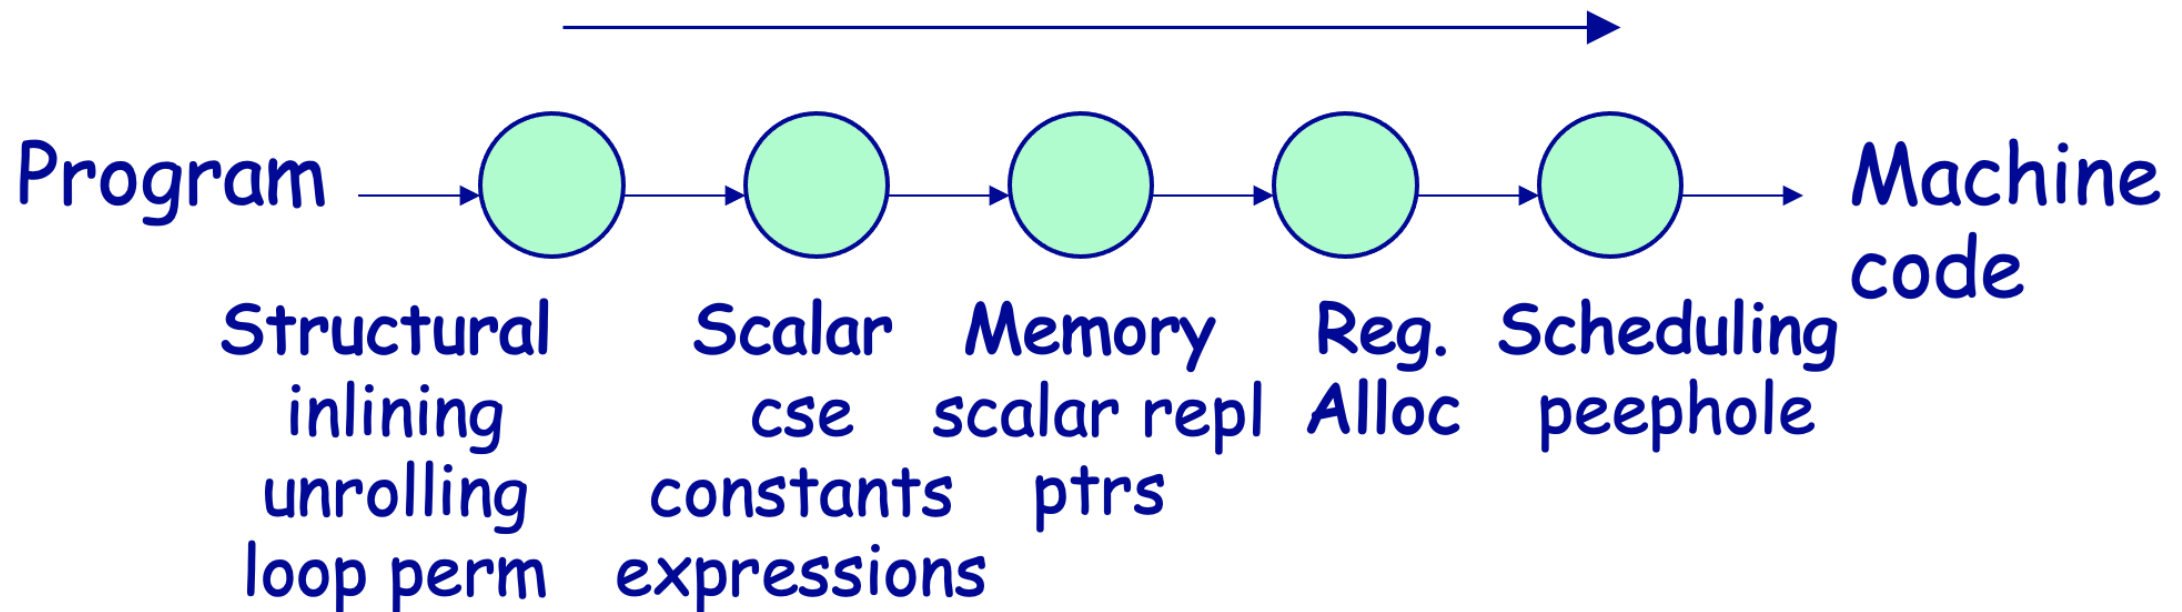
1. within a basic block (local)
2. within a method (global)
3. across methods (interprocedural)

Analysis

1. control flow graph - dominators, loops, etc.
2. dataflow analysis - flow of values
3. [static-single-assignment](#) – transform programs such that each variable has a unique definition
4. alias analysis - pointer memory usage
5. dependence analysis - array memory usage

Basic Compiler Structure

Higher to lower level
representations, analyses, & transformations



2. Static Analysis

1. Background

2. Static Analysis

- **Concept**
- Buffer Overrun
- Pitfalls of Privileges
- Untrusted Data

3. A Perfect Work

2.1 Concept

Static Analysis - Concept

- Examples: compiler optimizations, program verifiers
- Examine program text (no execution)
- Build a **model** of program state
 - An abstraction of the run-time state
- Reason over possible behaviors
 - E.g., “run” the program over the abstract state

Abstract interpretation

- Typically implemented via dataflow analysis
- Each program statement's transfer function indicates how it transforms state
- Example: What is the transfer function for `y = x++;`?

Selecting an abstract domain

$\langle x \text{ is odd}; y \text{ is odd} \rangle$

$y = x++;$

$\langle x \text{ is even}; y \text{ is odd} \rangle$

$\langle x \text{ is prime}; y \text{ is prime} \rangle$

$y = x++;$

$\langle x \text{ is anything}; y \text{ is prime} \rangle$

$\langle x = \{ 3, 5, 7 \}; y = \{ 9, 11, 13 \} \rangle$

$y = x++;$

$\langle x = \{ 4, 6, 8 \}; y = \{ 3, 5, 7 \} \rangle$

$\langle x=3, y=11 \rangle, \langle x=5, y=9 \rangle, \langle x=7, y=13 \rangle$

$y = x++;$

$\langle x=4, y=3 \rangle, \langle x=6, y=5 \rangle, \langle x=8, y=7 \rangle$

$\langle x_n = f(a_{n-1}, \dots, z_{n-1}); y_n = f(a_{n-1}, \dots, z_{n-1}) \rangle$

$y = x++;$

$\langle x_{n+1} = x_n + 1; y_{n+1} = x_n \rangle$

Research challenge: Choose good abstractions

- The abstraction determines the **expense** (in time and space)
- The abstraction determines the **accuracy** (what information is lost)
 - Less accurate results are poor for applications that require precision
 - Cannot conclude all true properties in the grammar

Static analysis: Characteristic

- Slow to analyze large models of state, so use abstraction
- Conservative: account for abstracted-away state
- Sound: (weak) properties are guaranteed to be true
 - Some static analyses are not sound

2.2 Buffer Overruns

1. Background

2. Static Analysis

- Concept
- **Buffer Overrun**
- Pitfalls of Privileges
- Untrusted Data

3. A Perfect Work

What Is a Buffer Overflow



Buffer Overflow

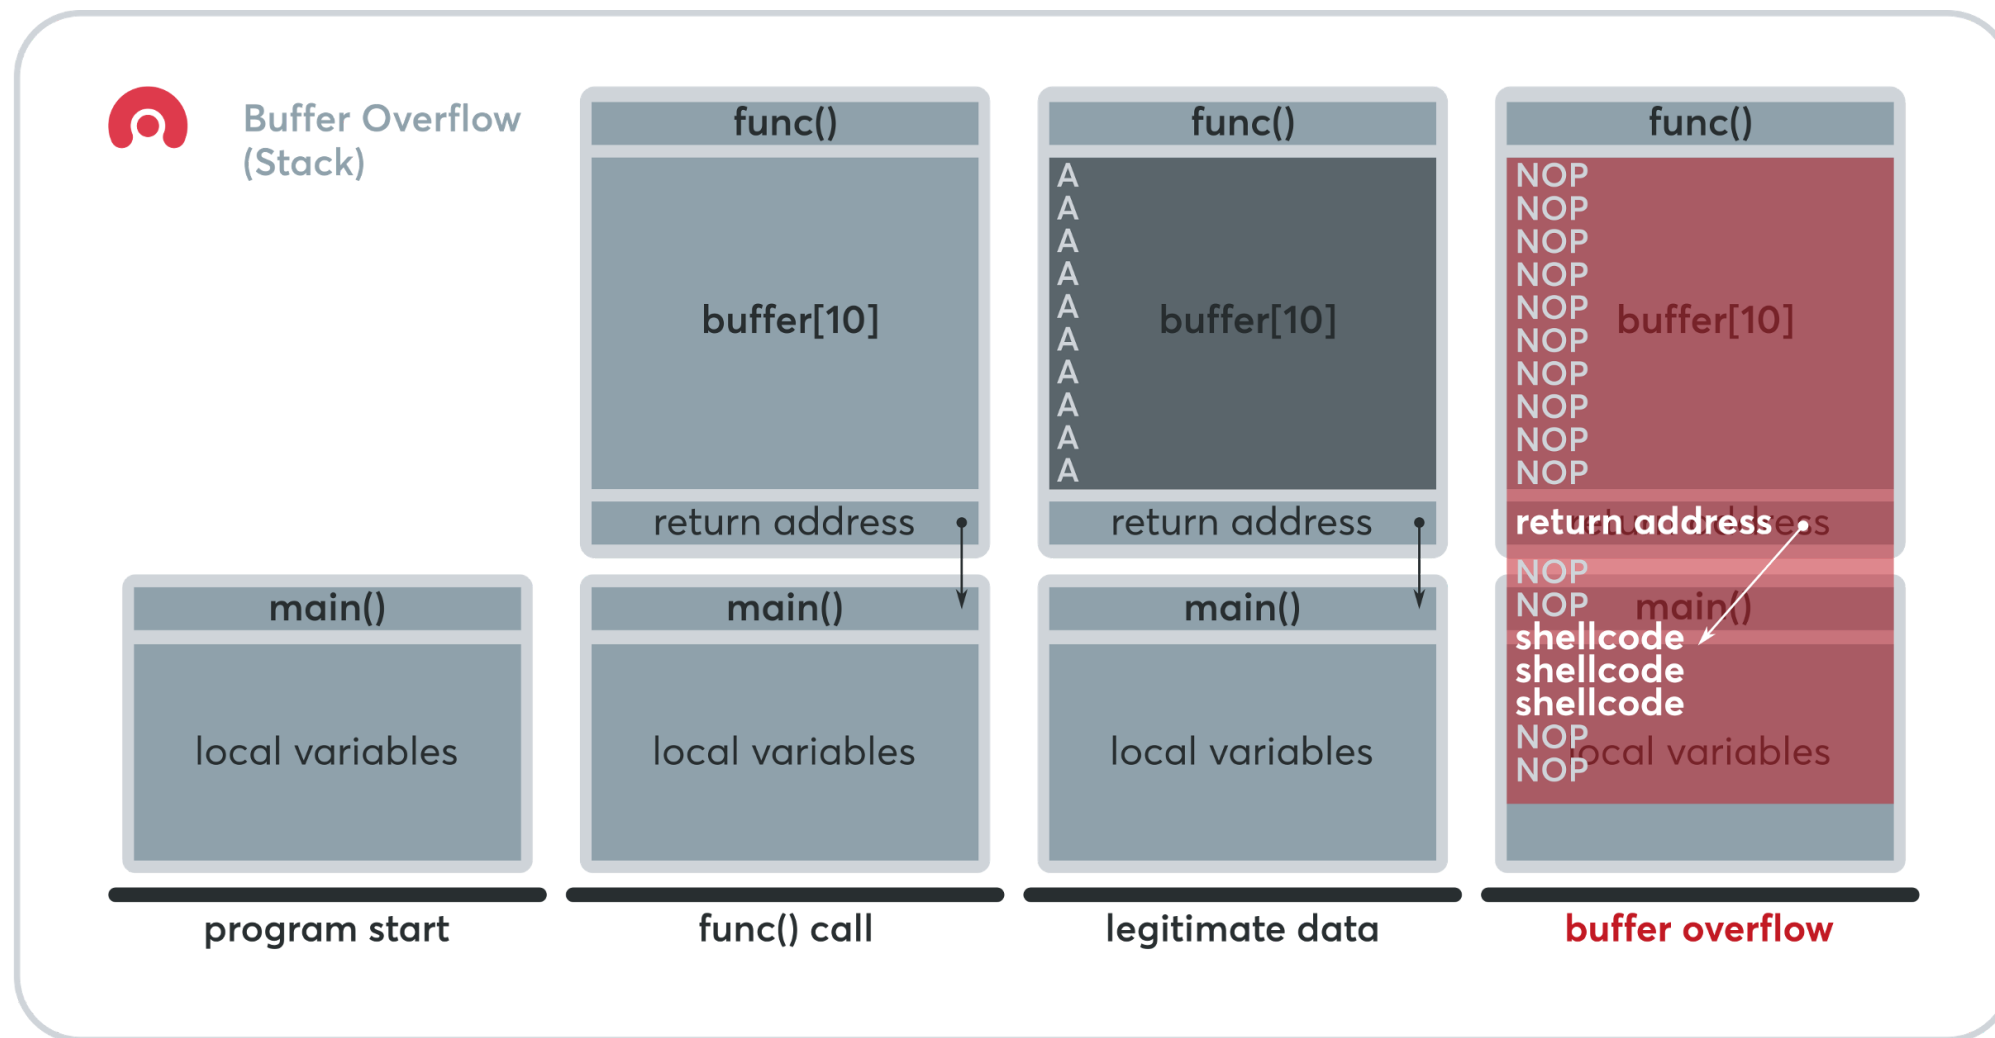
- An example bug

```
main(int argc, char *argv[]) {  
    func(argv[1]);  
}  
void func(char *v) {  
    char buffer[10];  
    strcpy(buffer, v);  
}
```

- Command line:

```
$ vulnprog AAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

Buffer Overflow



Static Detection of Overruns

- Introduce implicit variables:
 - `alloc(buf)` = # bytes allocated for buf
 - `len(buf)` = # bytes stored in buf
 - Safety condition: `len(buf) ≤ alloc(buf)`

Current Status

- Experimental results
 - Found new bugs in sendmail (30k LOC), others
 - Analysis is fast, but many false alarms (1/kLOC)
- Research challenges
 - Pointer analysis (support strong updates)
 - Integer analysis (infer linear relations, flow-sensitivity)
 - Soundness, scalability, real-world programs

2.3 Pitfalls of Privileges

1. Background

2. Static Analysis

- Concept
- Buffer Overrun
- **Pitfalls of Privileges**
- Untrusted Data

3. A Perfect Work

Pitfalls of Privileges

- Spot the bug:

```
setuid(0);  
rv = bind(...);  
if (rv < 0)  
    return rv;  
setuid(getuid());
```

Pitfalls of Privileges

- Spot the bug:

```
setuid(0);  
rv = bind(...);  
if (rv < 0)  
    return rv;  
seteuid(getuid());
```

enablePriv()

checkPriv()

Bug! Leaks privilege

disablePriv()

A Common Language

- Various interpretations are possible
 - C: enablePriv(p) lasts until next disablePriv(p)
 - Java: ... or until containing stack frame is popped
 - checkPriv(p) throws fatal error if p not enabled

Static Privilege Analysis

- Some problems in privilege analysis:
 - Privilege inference (auditing, bug-finding)
 - Find all privileges reaching a given program point
 - Enforcing privilege-safety (cleanliness of new code)
 - Verify statically that no `checkPriv()` operation can fail
 - ... or that program behaves same under C & Java styles

Future Directions

- Research challenges
 - Experimental studies on real programs
 - Handling data-directed privilege properties
 - Other access control models

2.4 Untrusted Data

1. Background

2. Static Analysis

- Concept
- Buffer Overrun
- Pitfalls of Privileges
- **Untrusted Data**

3. A Perfect Work

Manipulating Untrusted Data

- Spot the bug:

```
hp = gethostbyaddr(...);  
printf(hp->hp_hname);
```


Manipulating Untrusted Data

- Spot the bug:

```
hp = gethostbyaddr(...);  
printf(hp->hp_hname);
```

untrusted source of data

Bug! printf() trusts its first argument

Trust Analysis

- Security involves much mental “bookkeeping”
 - Problem: Help programmer keep track of which values can be trusted
- One approach: static taint analysis
 - Extend the C type system
 - Qualified types express annotations: e.g., `tainted char *` is an untrusted string
 - Typechecking enforces safe usage
 - Type inference reduces annotation burden

A Tiny Example

```
void printf(untainted char *, ...);  
tainted char * read_from_network(void);  
  
tainted char *s = read_from_network();  
printf(s);
```

After Type Inference...

```
void printf(untainted char *, ...);  
tainted char * read_from_network(void);
```

an inferred type

```
tainted char *s = read_from_network();  
printf(s);
```

Doesn't type-check!
Indicates vulnerability

... where untainted $T \leq$ tainted T

Current Status

- Experimental results
 - Successful on real programs
 - Able to find many previously-known format string bugs
 - Cost: 10-15 minutes per application
 - Type theory seems useful for security engineering
- Research challenges
 - Richer theory to support real programming idioms
 - More broadly-applicable discipline of good coding
 - Finer-grained notions of trust

Concluding Remarks

- Static analysis can help secure our software
 - Buffer overruns, privilege bugs, format string bugs
 - Hits a sweet spot: cheap and proactive
- Security as a source of interesting problems?
 - Motivations for better pointer, integer analysis
 - New problems: privilege analysis, trust analysis

3. A Perfect Work

- Improving Integer Security for Systems with KINT
 - <https://www.usenix.org/conference/osdi12/technical-sessions/presentation/wang>
- Towards Optimization-Safe Systems: Analyzing the Impact of Undefined Behavior
 - <http://sigops.org/sosp/sosp13/program.html>

References

1. [Lint, a C Program Checker](#)
2. [Towards Optimization-Safe Systems: Analyzing the Impact of Undefined Behavior](#)
3. [Improving Integer Security for Systems with KINT](#)
4. [Static Analysis and Software Assurance](#)
5. [Static and dynamic analysis: synergy and duality](#)
6. [compiler](#)
7. [Static/Dynamic Analysis Tools](#)
8. [C Programming Tools](#)

谢谢！